

Overview of Quasi-Experimental Designs

Advanced Social Epidemiology PhD Course

Sam Harper

University of Copenhagen
2021-10-11 to 2021-10-15

4. Quasi-Experiments

4.1. Motivation

4.2. Randomization and Observation

4.2 Quasi-Experimental Designs

4. Quasi-Experiments

4.1. Motivation

4.2. Randomization and Observation

4.2 Quasi-Experimental Designs

Stylized "forms" of questions asked in social epidemiology

What question do most studies in social epidemiology answer?

- Do individuals who are disadvantaged with respect to social position have worse health than those who are advantaged?

Other kinds of questions that could be asked:

- **Would** individuals who are disadvantaged with respect to social position have better health **if they were to become advantaged?**
- **Would** individuals who are advantaged with respect to social position have worse health **if they were to become disadvantaged?**

These are **causal** questions.

"Normal" etiological science in social epidemiology

1. Follow-up of individuals in different social groups for various health outcomes (incidence, mortality, risk factors)
2. Adjustment for various confounders/mediators (are inequalities "explained" by....A, B, C?).
 - "Our results demonstrate that"...we should:
 - *raise* education levels
 - *increase* economic assistance to the poor
 - *remove* noxious exposures from the environment
 - *reduce* psychosocial workplace hazards
 - *eliminate* hierarchies, and the like.
 - These statements are based on making **causal** inferences.

What's the problem?

- We are mainly (though not exclusively) interested in causal effects.
- We want to know:
 - Should we intervene to reduce exposure to X ? or
 - Did the program work? If so, for whom? If not, why not?, or
 - If we implement the program elsewhere, should we expect the same result?
- These questions involve counterfactuals about what would happen **if** we intervened to do something.
- These are causal questions.

Causation, Association, and Confounding

- **Causal effect:** Do individuals randomly **assigned** (i.e., SET) to treatment have better outcomes?

$$E(Y|SET[Treated]) - E(Y|SET[Untreated])$$

- **Association:** Do individuals who **happen to be** treated have better outcomes?

$$E(Y|Treated) - E(Y|Untreated)$$

- **Confounding:**

$$E(Y|SET[Treated]) - E(Y|SET[Untreated]) \neq E(Y|Treated) - E(Y|Untreated)$$

4. Quasi-Experiments

4.1. Motivation

4.2. Randomization and Observation

4.2 Quasi-Experimental Designs

Randomized Trials vs. Observational Studies

RCTs, Defined

RCTs involve:

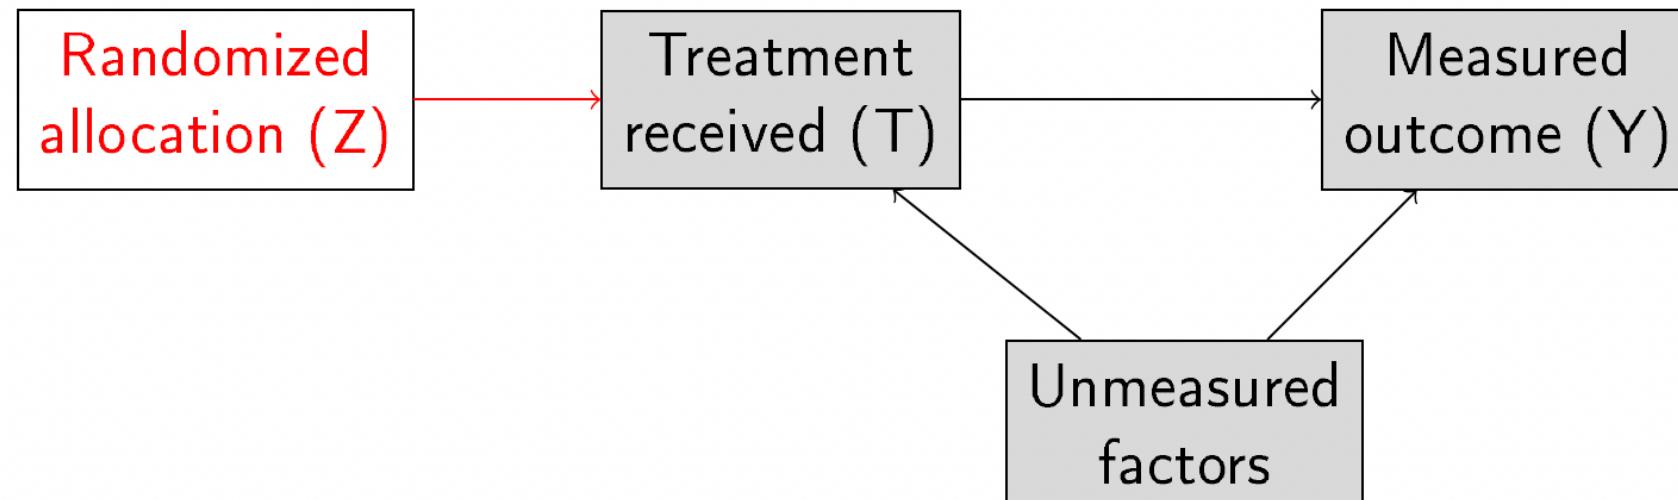
1. comparing treated and control groups;
2. the treatment assignment is random;
3. investigator does the randomizing.

In an RCT, treatment/exposure is **assigned** by the investigator

- In observational studies, exposed/unexposed groups **exist** in the source population and are selected by the investigator.
- Good natural experiments do (1) and (2), but not (3).
- Because there is no control over assignment, the credibility of natural experiments hinges on how good "as-if random" approximates (2).

Strength of randomized treatment allocation

- Recall that randomization means that we can generally estimate the causal effect without bias.
- Randomization guarantees exchangeability on measured and unmeasured factors.



Randomize if you can

- Randomization leads to:
 - balance on measured factors.
 - balance on unmeasured factors.
- Unmeasured factors cannot bias the estimate of the exposure effect.
- Example from Home Injury Prevention Intervention cluster RCT
- What do you notice about Table 1?

(Keall et al. 2015)

	Treatment group (n=950)	Control group (n=898)
Female sex	541 (57%)	501 (56%)
Indigenous Māori	88 (9%)	86 (10%)
Mean (SD) age (years)*	45 (28.0)	43 (28.1)
Age range (years)	0–94	0–92
0–9	175 (18%)	187 (21%)
10–19	89 (9%)	82 (9%)
20–29	34 (4%)	37 (4%)
30–39	116 (12%)	112 (12%)
40–49	90 (9%)	96 (11%)
50–59	65 (7%)	51 (6%)
60–69	132 (14%)	105 (12%)
≥70	249 (26%)	228 (25%)
Number of injuries at home, excluding falls, in past year (per person)†	122 (0.129)	103 (0.115)
Number of fall injuries at home in past year (per person)‡	87 (0.092)	61 (0.068)
Number of specific injuries in past year (per person)§	23 (0.024)	24 (0.027)

Data are number of individual occupants (%), unless otherwise indicated.
*At Aug 3, 2010. †Injuries arising in the home during the 365-day period before the intervention date, obtained from matched insurance claim data. ‡Slips, trips, or fall injuries in the home during the 365-day period before the intervention date. §Injuries most specific to the package of home modifications, arising in the home during the 365-day period before the intervention date.

Table 1: Characteristics of individual occupants at baseline

What's the problem?

We are mainly (though not exclusively) interested in causal effects.

Randomization is generally great for answering whether treatment assignment Z affects Y .

- treatment assignment (Z) is independent of potential outcomes and all measured and **unmeasured** pre-treatment variables.
- Effect of Z on Y is unconfounded

RCTs have serious limitations.

- Non-compliance.
- Attrition.
- Spillovers.
- Blinding (esp. in clinical trials).

Problem of Social Exposures

- Many social exposures cannot be randomized by investigators:
 - Unethical (poverty, parental social class, job loss)
 - Impossible (ethnic background, place of birth)
 - Expensive (neighborhood environments)
- Some exposures are hypothesized to have long latency periods (many years before outcomes are observable).
- Effects may be produced by complex, intermediate pathways.
- We need alternatives to RCTs.

How to interpret statistical associations of health inequality?

We have lots of statistical associations between social exposures and health.

$$X - - - Y$$

Some possible situations *consistent* with statistical associations:

1. Causal $X \rightarrow Y$
2. Heterogeneity $X_a \ Y_a$ vs. $X_b \rightarrow Y_b$
3. Reverse causation $Y \rightarrow X$
4. Confounding $X \leftarrow C \rightarrow Y$
5. Selection bias $X \rightarrow S \leftarrow Y$

Unmeasured confounding is a **serious** challenge

- We often compare socially advantaged and disadvantaged on health.
- Key problem: people choose/end up in treated or untreated group for reasons that are difficult to measure and that may be correlated with their outcomes.
- So...**adjust**.
 - Measure and adjust (regression) for C confounding factors.
 - Conditional on C , we are supposed to believe assignment is "as good as random" = causal.

Key issue is credibility

- If we have a good design and assume that we have measured all of the confounders, then regression can give us exactly what we want: an estimate of the causal effect of exposure to T .
- Core issue: How credible is this assumption?



"Now, keep in mind that these numbers are only as accurate as the fictitious data, ludicrous assumptions and wishful thinking they're based upon!"

SEP and CVD in Australia. Many low p-values

Table 1 Characteristics of 38 355 subjects in the Melbourne Collaborative Cohort Study at baseline (1990–1994)

		Highest level of education attained				
		Completed tertiary* n = 8588	Completed secondary† n = 7882	Some secondary‡ n = 14543	Primary only§ n = 7342	p Value for trend¶
Male	n (%)	4025 (47%)	3776 (48%)	4680 (32%)	2780 (38%)	<0.001
Female	n (%)	4563 (53%)	4106 (52%)	9863 (68%)	4562 (62%)	<0.001
Age (years)	(Mean, SD)	51.6 (8.4)	54.5 (8.8)	55.7 (8.5)	57.8 (7.1)	<0.001
Country of birth, n (%)	Australia, New Zealand or northern Europe (n=28 835)	8263 (96%)	6814 (86%)	12696 (87%)	1062 (14%)	<0.001
	Southern Europe (n=9520)	325 (4%)	1068 (14%)	1847 (13%)	6280 (86%)	<0.001
Behavioural risk factors						
Current smoker	n (%)	574 (7%)	960 (12%)	1828 (13%)	947 (13%)	<0.001
Vegetable intake (times/day)	Mean (SD)	5.7 (3)	5.3 (3)	5.2 (3)	5.8 (4)	1.000
Fruit intake (times/day)	Mean (SD)	4.4 (3)	4.0 (3)	3.9 (3)	4.7 (4)	0.007
Saturated fat intake (g/day)	Mean (SD)	35.0 (15)	34.3 (16)	33.7 (16)	30.3 (18)	<0.001
Current drinker	n (%)	7061 (82%)	5883 (75%)	9397 (65%)	3666 (50%)	<0.001
Alcohol intake, current drinkers (g/day)	Median (IQR)	14 (5,26)	13 (4,26)	10 (3, 23)	15 (4,30)	<0.001
Physical activity (% inactive)	n (%)	1224 (14%)	1520 (19%)	3238 (22%)	2546 (35%)	<0.001
Social connection						
Living alone	n (%)	1514 (18%)	1274 (16%)	2250 (15%)	498 (7%)	<0.001

Is observational credibility is getting harder to sell?

Another example: Does breastfeeding increase child IQ?

Several observational studies show higher IQs for breastfed children.

"The authors of this and other studies claim to find effects of breastfeeding because even once they adjust for the differences they see across women, the effects persist. But this assumes that the adjustments they do are able to remove all of the differences across women. This is extremely unlikely to be the case."

"I would argue that in the case of breastfeeding, this issue is impossible to ignore and therefore **any study that simply compares breastfed to formula-fed infants is deeply flawed**. That doesn't mean the results from such studies are necessarily wrong, just that we can't learn much from them."

Oster (2015)

Why we worry about observational studies

Recent evaluation of "Workplace Wellness" program in US state of Illinois

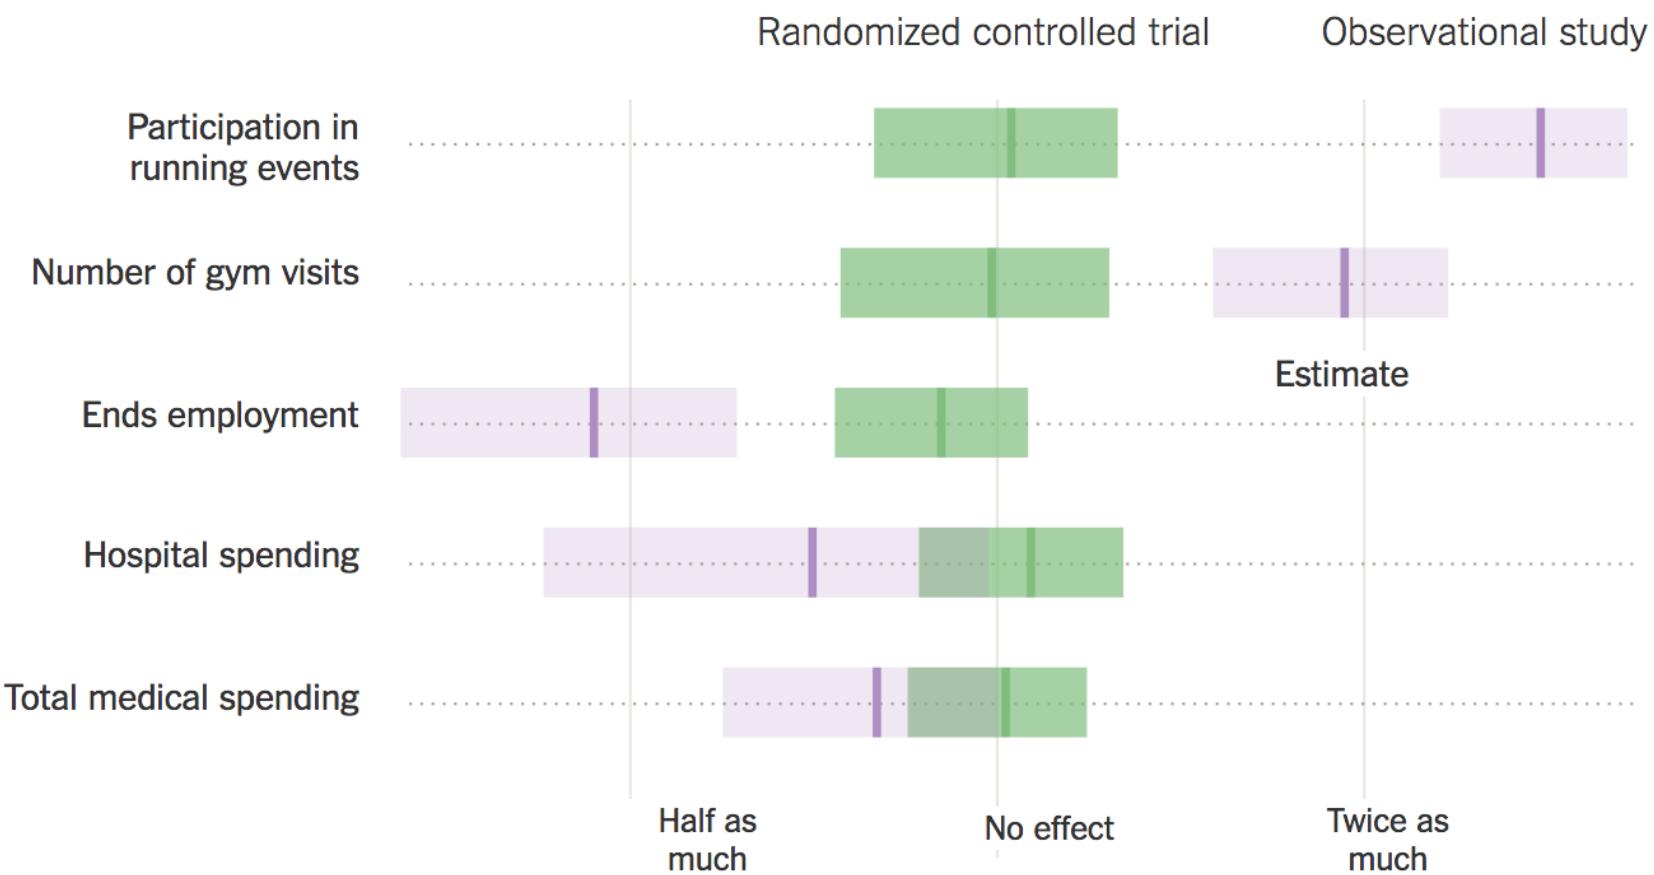
Treatment: biometric health screening; online health risk assessment, access to a wide variety of wellness activities (e.g., smoking cessation, stress management, and recreational classes).

Randomized evaluation:

- 3,300 individuals assigned treated group.
- 1,534 assigned to control (could not access the program).

Also analyzed as an observational study comparing "participants" vs. non-participants in treated group.

How the Illinois Wellness Program Affected ...



Carroll, New York Times, Aug 6, 2018.

4. Quasi-Experiments

4.1. Motivation

4.2. Randomization and Observation

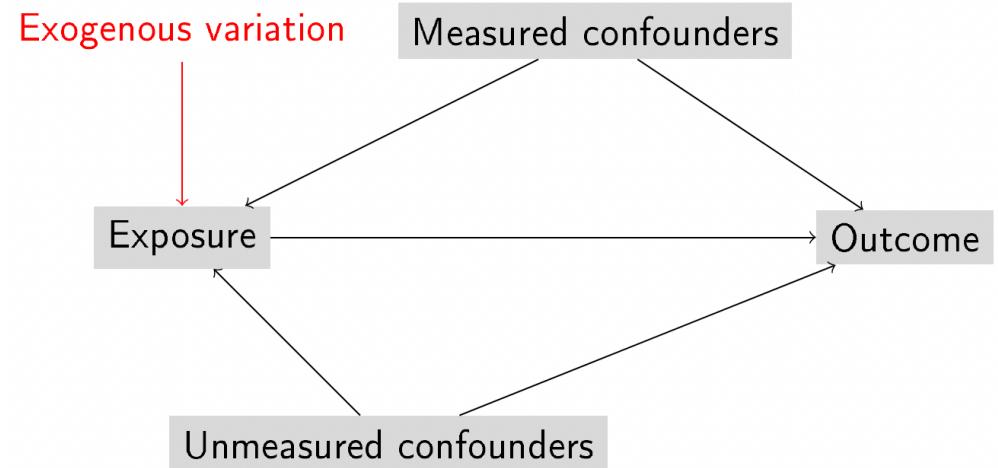
4.2 Quasi-Experimental Designs

How can quasi-experiments help?

- Quasi-experiments aim to mimic RCTs.
- "Accidents of chance" that create:
 1. Comparable treated and control units
 2. Random or "as-if" random assignment to treatment.
- Control for (some) sources of bias that cannot be adequately controlled using regression adjustment.
- More credible designs also help us to understand the relevance of other factors that may be implicated in generating inequalities.

Selection on "observables" and "unobservables"

- Observables: Things you measured or can measure
- Unobservables: Things you can't measure (e.g., innate abilities)
- Exogenous variation: predicts exposure but (**we assume**) **not** associated with anything else [mimicking random assignment].



Strategies based on observables and unobservables

- Most observational study designs control for *measured* factors using:
 - Stratification
 - Regression adjustment
 - Matching (propensity scores, etc.)
- Quasi-experimental strategies **aim** to control for some *unmeasured* factors using:
 - Interrupted time series (ITS)
 - Difference-in-differences (DD)
 - Synthetic controls (SC)
 - Instrumental variables (IV)
 - Regression discontinuity (RD)

Some *potential* sources of natural experiments

- Law changes
- Eligibility for social programs (roll-outs)
- Lotteries
- Genes
- Weather shocks (rainfall, disasters)
- Arbitrary policy or clinical guidelines (thresholds)
- Business / factory closures
- Historical legacies (physical environment)
- Seasonality

Difference-in-Differences

Difference-in-Differences: Basic Idea

In the simplest DD setting, outcomes are observed for units in two groups and in two time periods.

Treated:

- only units in one of the two groups are exposed to a treatment, in the second time period.

Control:

- Never observed to be exposed to the treatment.

Difference-in-Differences: Basic Idea

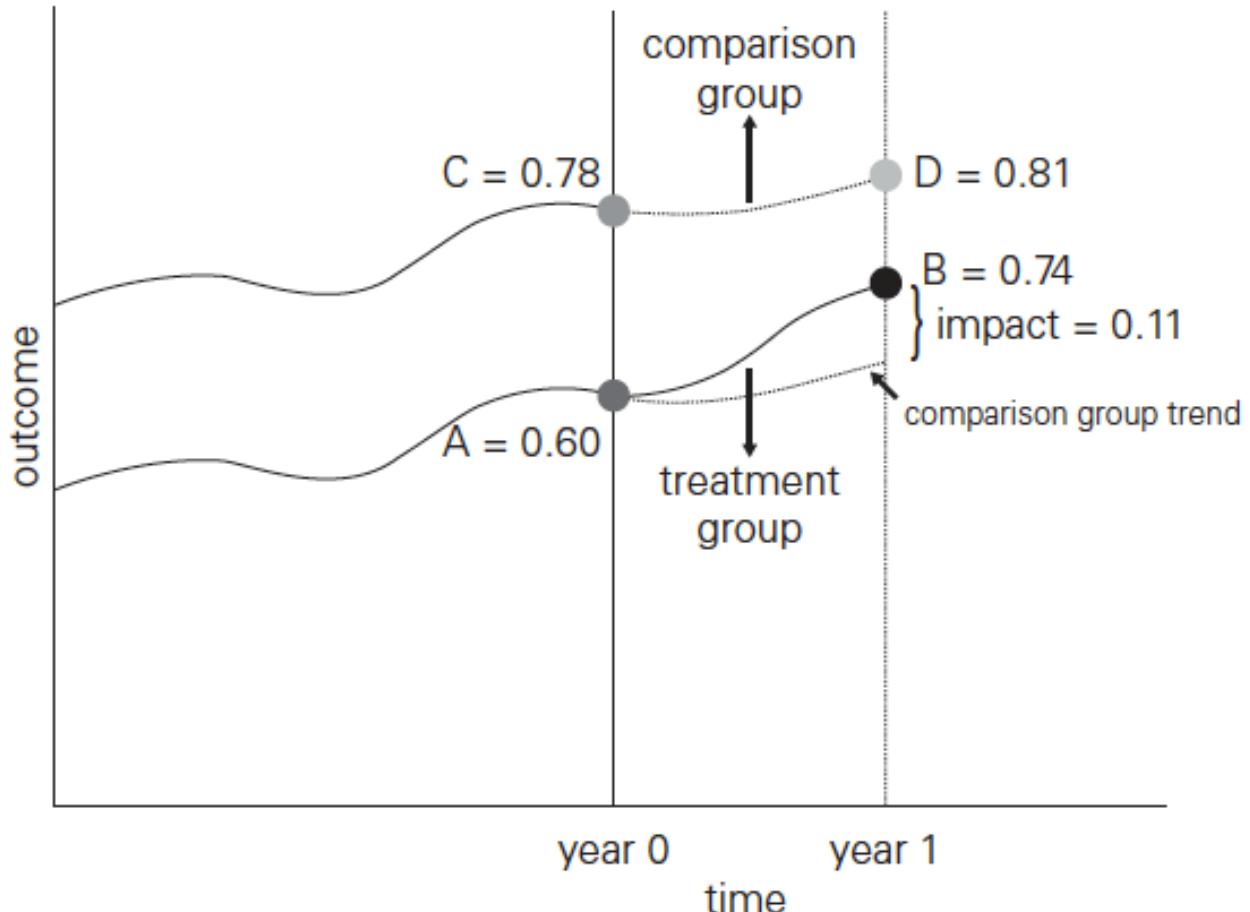
The average **change** over time in the non-exposed (control) group is **subtracted** from the gain over time in the exposed (treatment) group.

These are our two 'differences'.

Double differencing removes biases in second period comparisons between the treatment and control group that could be the result from

- permanent differences between those groups
- secular trends affecting both groups.

Visual Intuition of DD



Difference-in-Differences without Regression

DD is just differences in means! Let $\mu_{it} = E(Y_{it})$

- $i = 0$ is control group, $i = 1$ is treatment.
- $t = 0$ is pre-period, $t = 1$ is post-period.
- One 'difference' estimate of causal effect is: $\mu_{11} - \mu_{10}$ (pre-post in treated)
- Differences-in-Differences estimate of causal effect is: $(\mu_{11} - \mu_{10}) - (\mu_{01} - \mu_{00})$

Area	Before	After	Difference (A - B)
Treated	135	100	-35
Control	80	60	-20
T - C	55	40	-15

A social epidemiology example

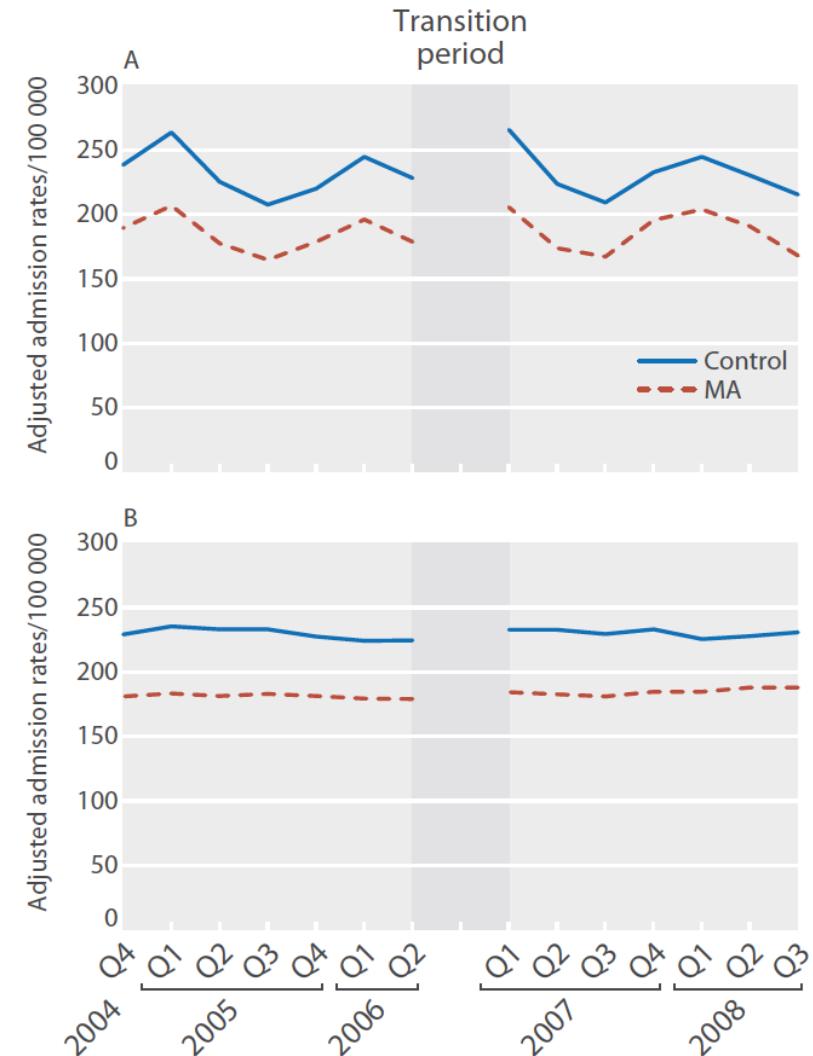
Effect of Massachusetts healthcare reform on racial and ethnic disparities in admissions to hospital for ambulatory care sensitive conditions: retrospective analysis of hospital episode statistics

Danny McCormick,¹ Amresh D Hanchate,^{2,3} Karen E Lasser,³ Meredith G Manze,³ Mengyun Lin,³ Chieh Chu,³ Nancy R Kressin^{2,3}

- Evaluated impact of MA reform on inequalities in hospital admissions.
- Compared MA to nearby states: NY, NJ, PA.
- Intervention "worked": % uninsured halved (12% to 6%) from 2004-06 to 2008-09.

Evaluating pre-intervention trends

- Adds credibility to assumption that post-intervention trends **would have been similar** in the absence of the intervention.
- "Null" results help focus on alternative mechanisms linking disadvantage to hospital admissions.



Synthetic Controls

Synthetic control methods

- Inference from comparative case studies is limited if we cannot identify a control to represent the counterfactual scenario.
- Abadie and Gardeazabel [@Abadie:2003aa] pioneered the synthetic control method to examine the economic impact of terrorism in the Basque country, using other Spanish regions as control groups.
- The synthetic control method uses a data driven approach to compare the trend of an outcome in a treated unit with the trend in a synthetic composite area (the "synthetic control").
- The method is gaining popularity in social sciences (i.e., political science), but hasn't taken root in epidemiology as of yet.



Contents lists available at [ScienceDirect](#)

Journal of Health Economics

journal homepage: www.elsevier.com/locate/econbase



The effects of paid maternity leave: Evidence from Temporary Disability Insurance



Jenna Stearns^{*,1}

University of California, Santa Barbara, United States

ARTICLE INFO

Article history:

Received 30 October 2014

Received in revised form 13 April 2015

Accepted 28 April 2015

Available online 7 May 2015

Keywords:

Temporary Disability Insurance

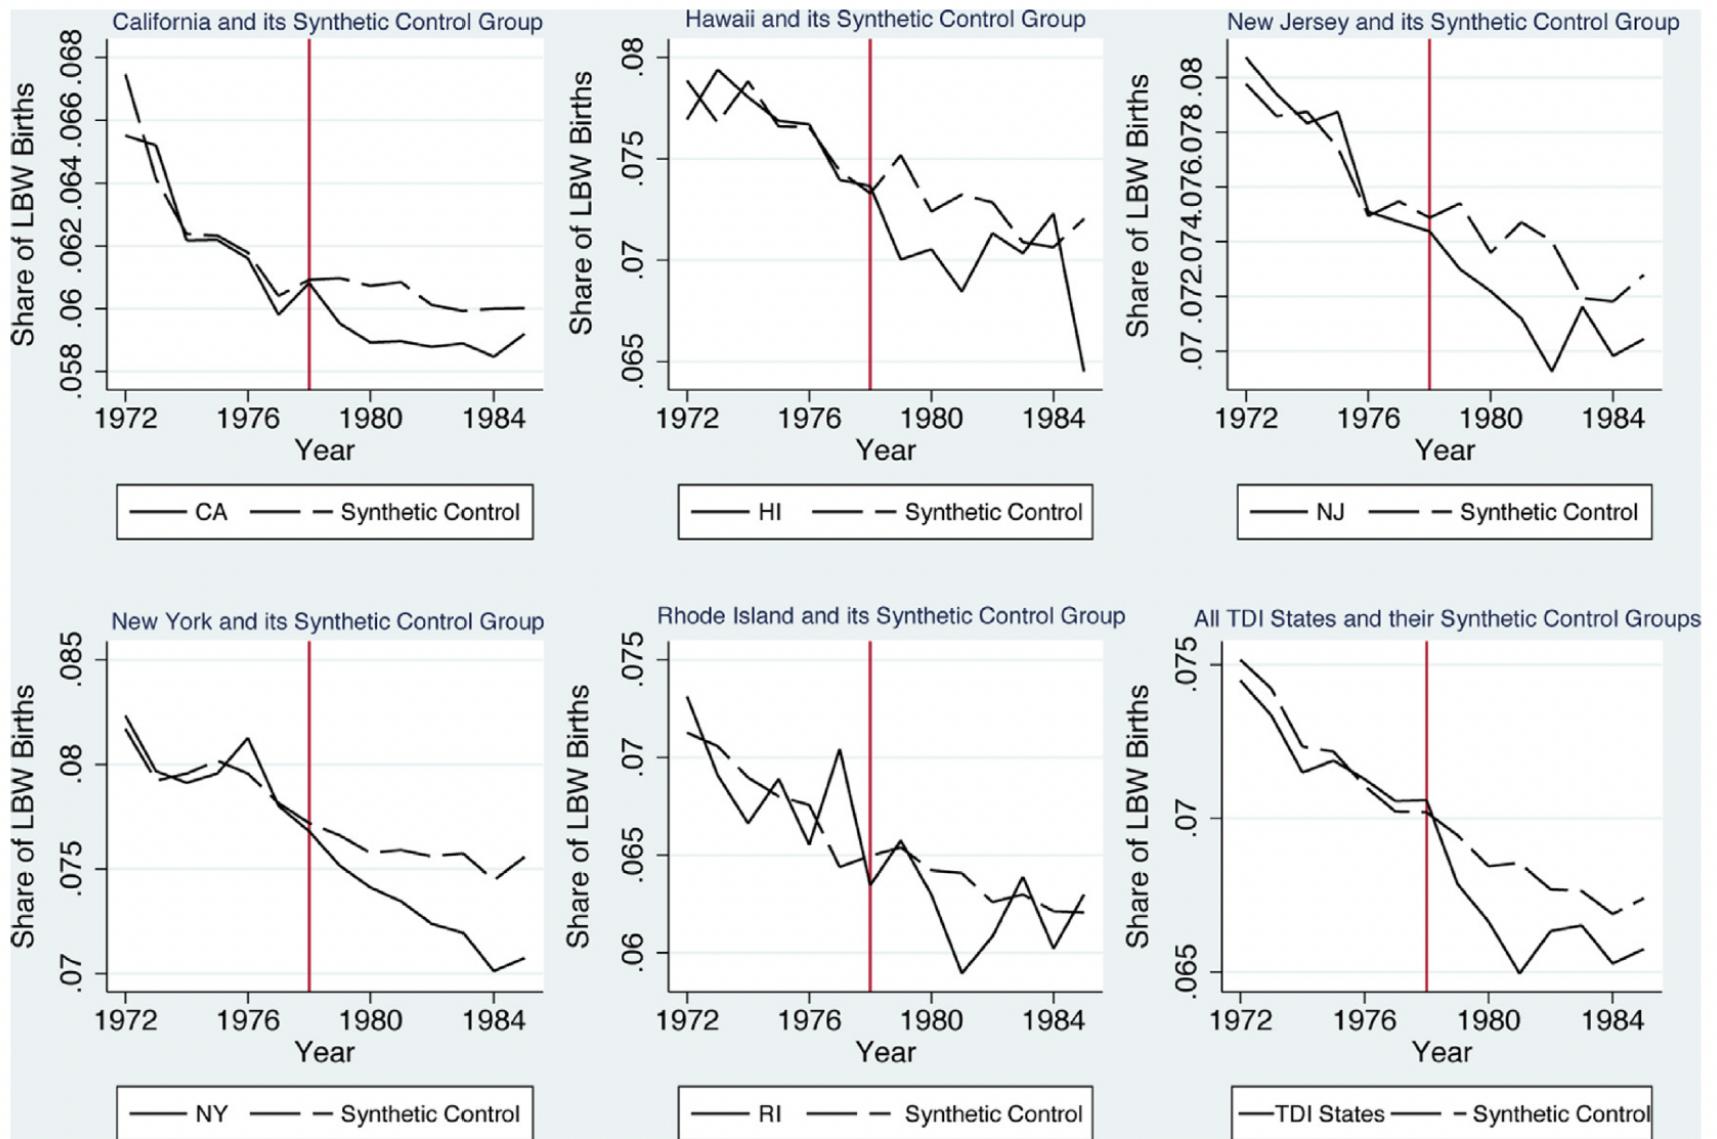
Maternity leave

Infant health

ABSTRACT

This paper investigates the effects of a large-scale paid maternity leave program on birth outcomes in the United States. In 1978, states with Temporary Disability Insurance (TDI) programs were required to start providing wage replacement benefits to pregnant women, substantially increasing access to antenatal and postnatal paid leave for working mothers. Using natality data, I find that TDI paid maternity leave reduces the share of low birth weight births by 3.2 percent, and the estimated treatment-on-the-treated effect is over 10 percent. It also decreases the likelihood of early term birth by 6.6 percent. Paid maternity leave has particularly large impacts on the children of unmarried and black mothers.

© 2015 Elsevier B.V. All rights reserved.



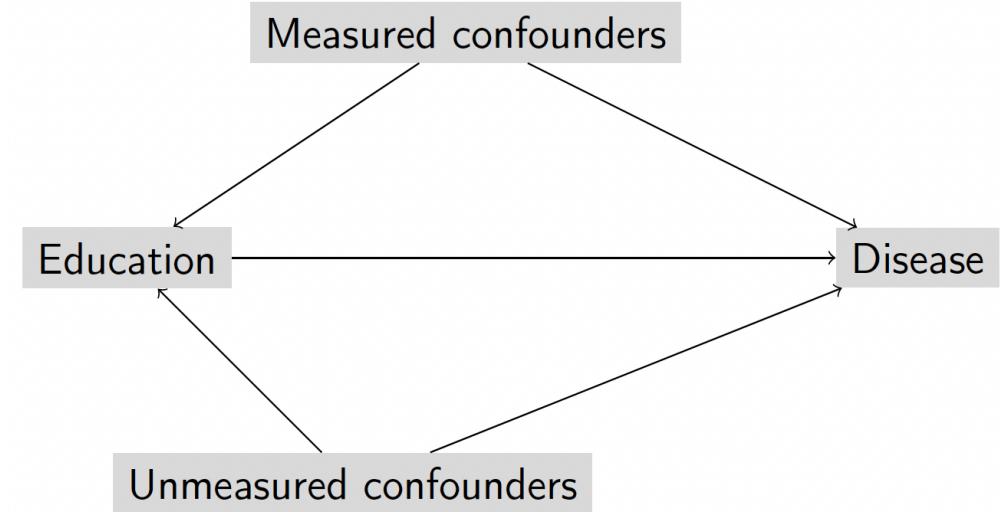
Stearns (2015)

Fig. 1. Average share of low birth weight births.

Instrumental Variables

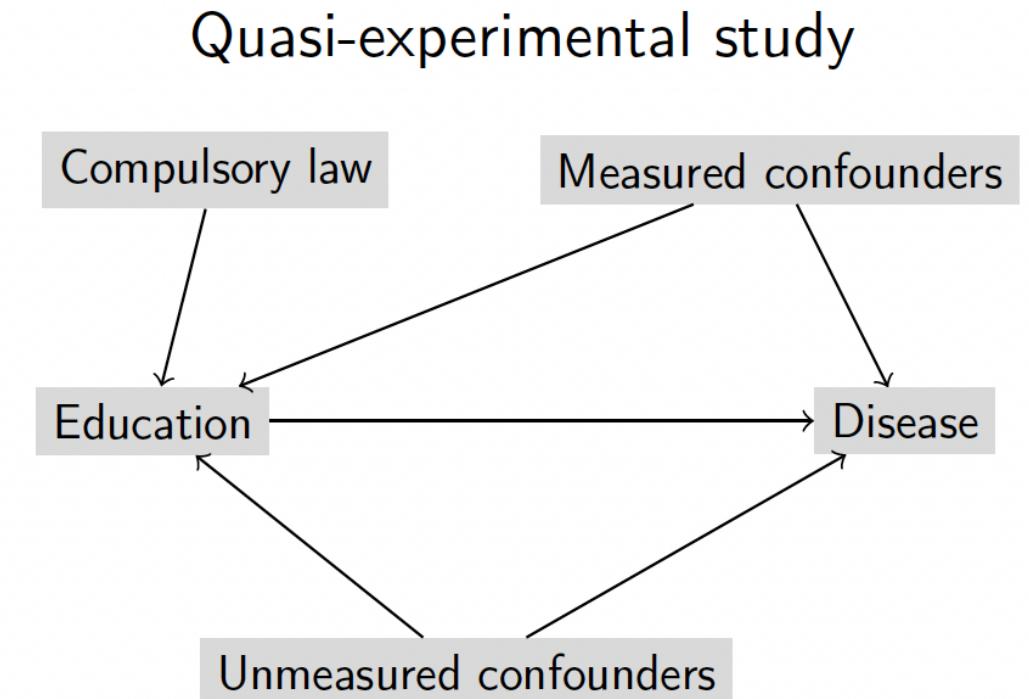
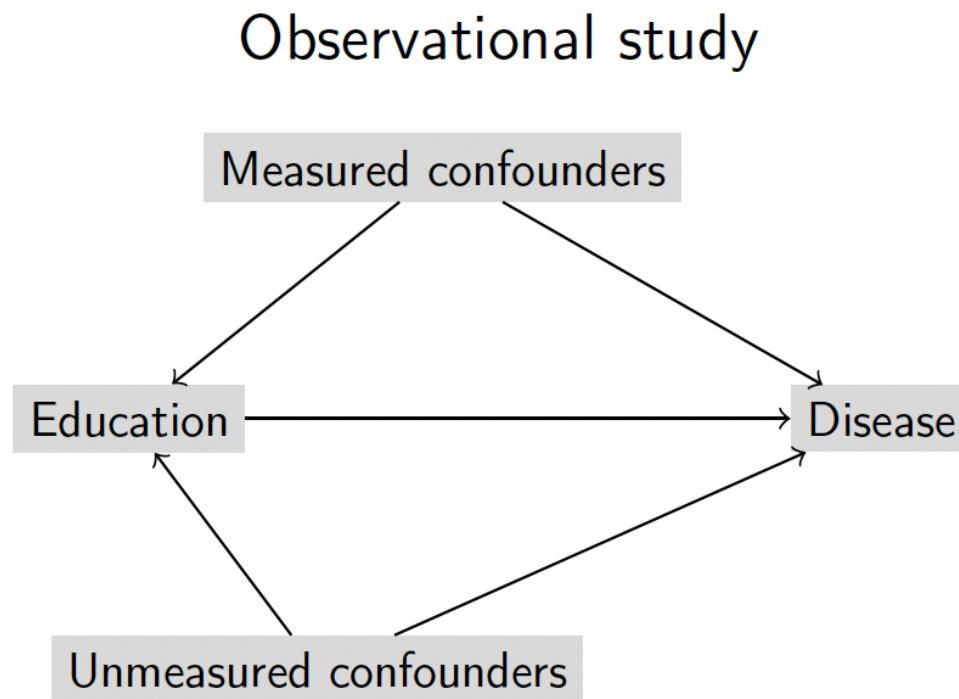
Challenge of conventional observation study (again)

- WHO: "Educational attainment is linked to improved health outcomes."
- But what about unmeasured confounding? Unmeasured factors such as personality traits, cognitive ability, etc. may be predictive of both education and disease.
- Failure to measure such factors will falsely attribute their effects to education.



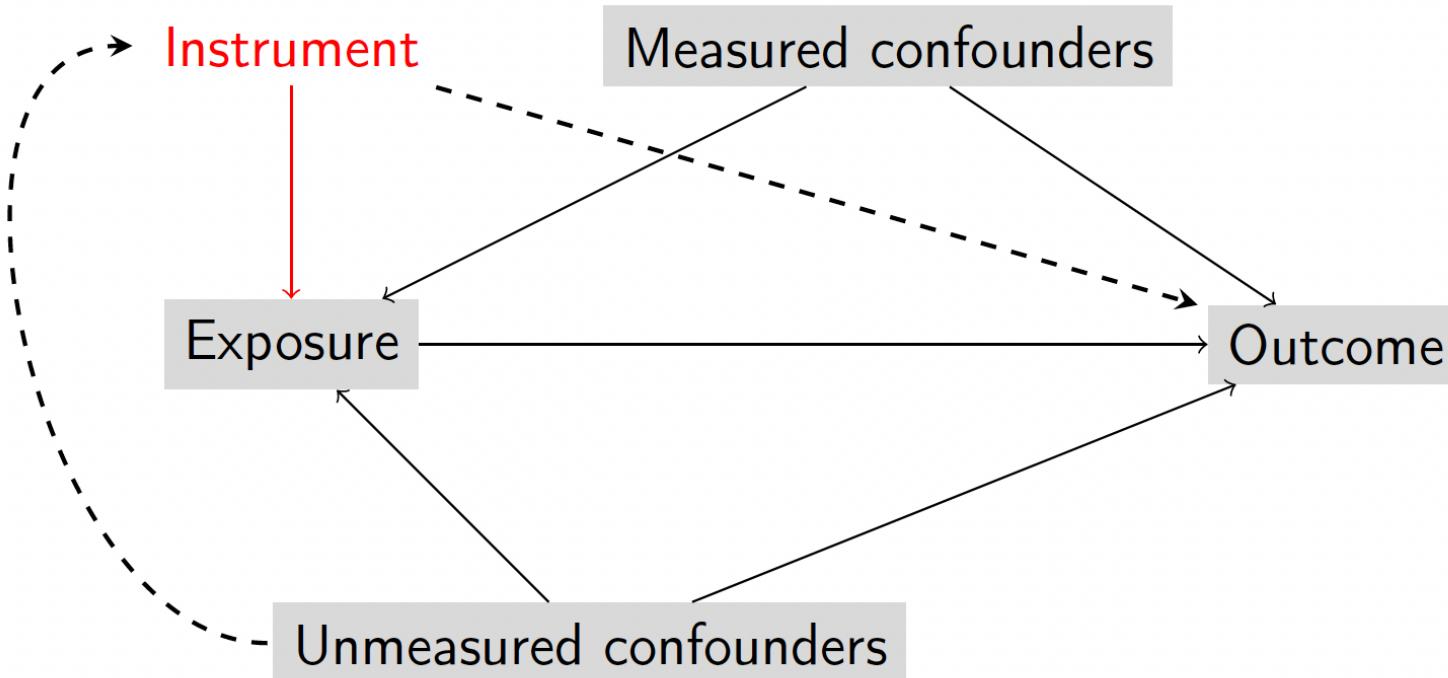
Possible solution: Quasi-experiment

"Instrumental variable": predicts education but **not** associated with anything else [mimicking random assignment].



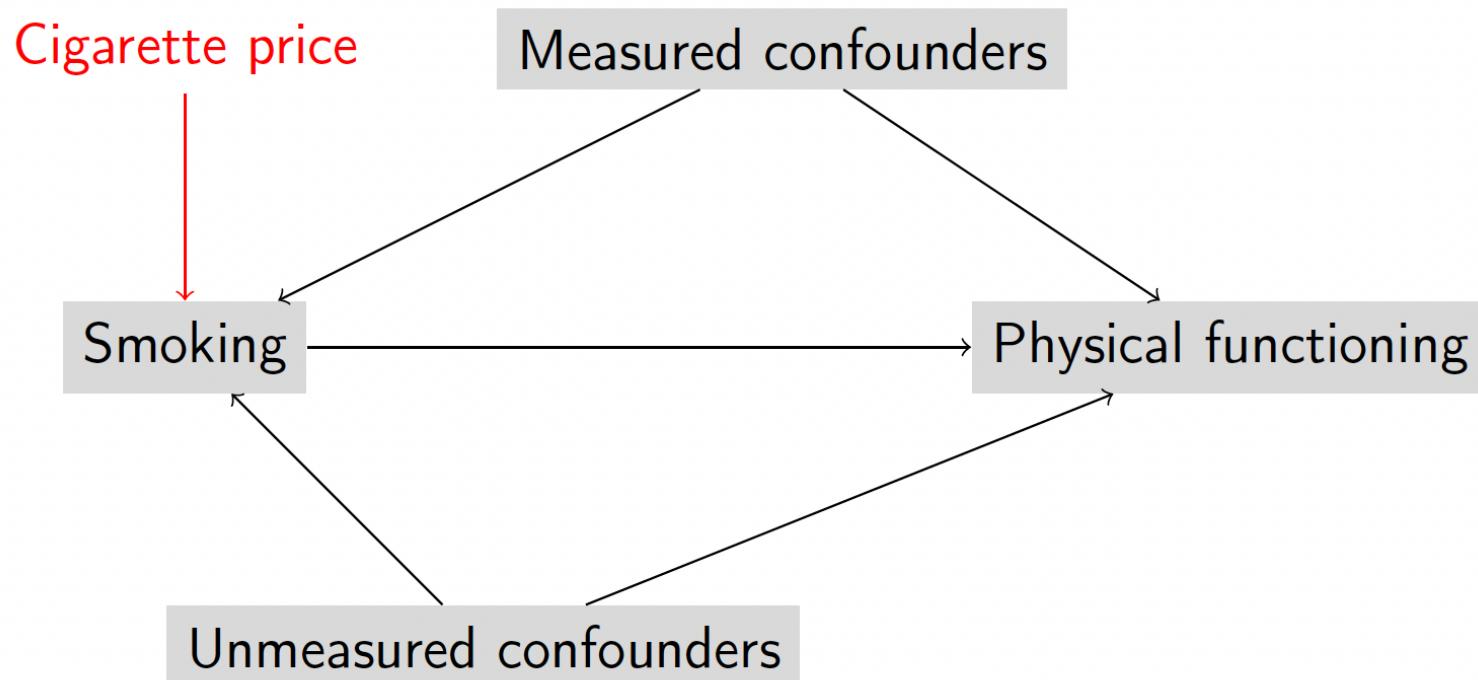
Non-randomized instrument creates additional issues

- In an RCT we know the treatment assignment is not associated directly with the outcome or with other unmeasured common causes.
- This assumption is less credible when our "instrument" is non-randomized.



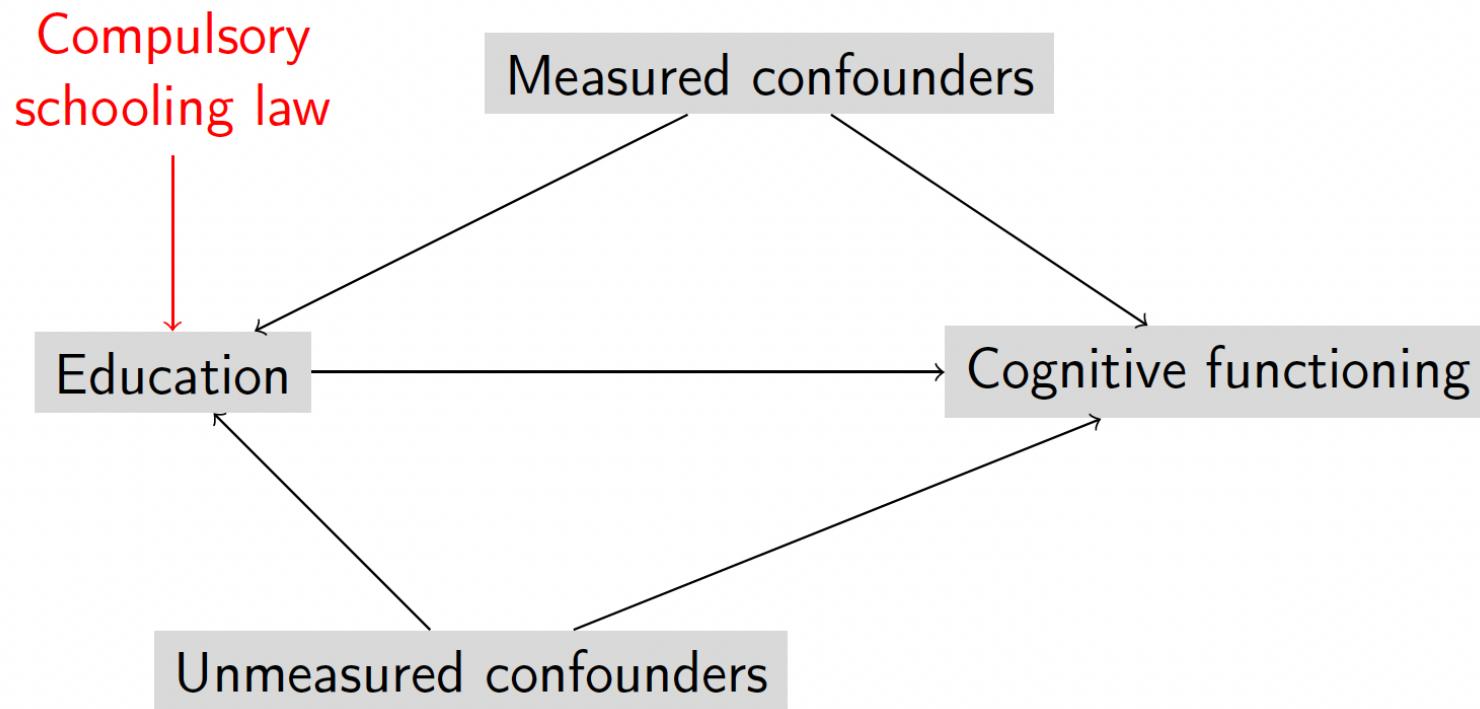
Non-randomized examples of IV: Policies

- Does smoking affect physical functioning ?
- **Instrument:** changes in cigarette prices [mimicking random assignment].



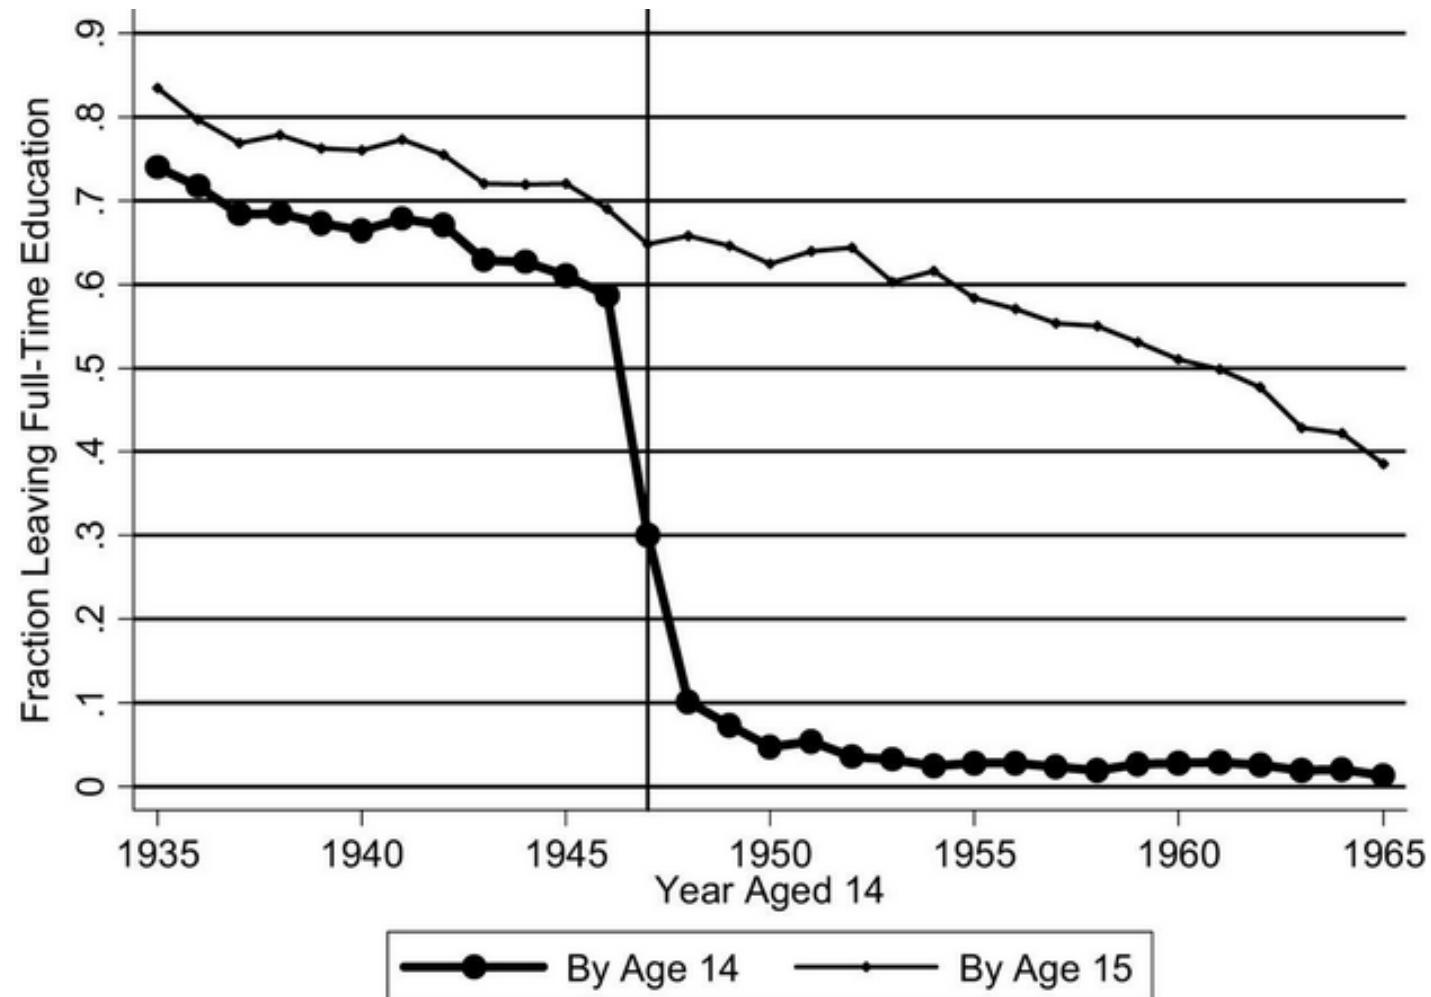
Non-randomized examples of IV: Policies

- Does education affect cognitive functioning?
- **Instrument:** changes in compulsory schooling laws [mimicking random assignment].



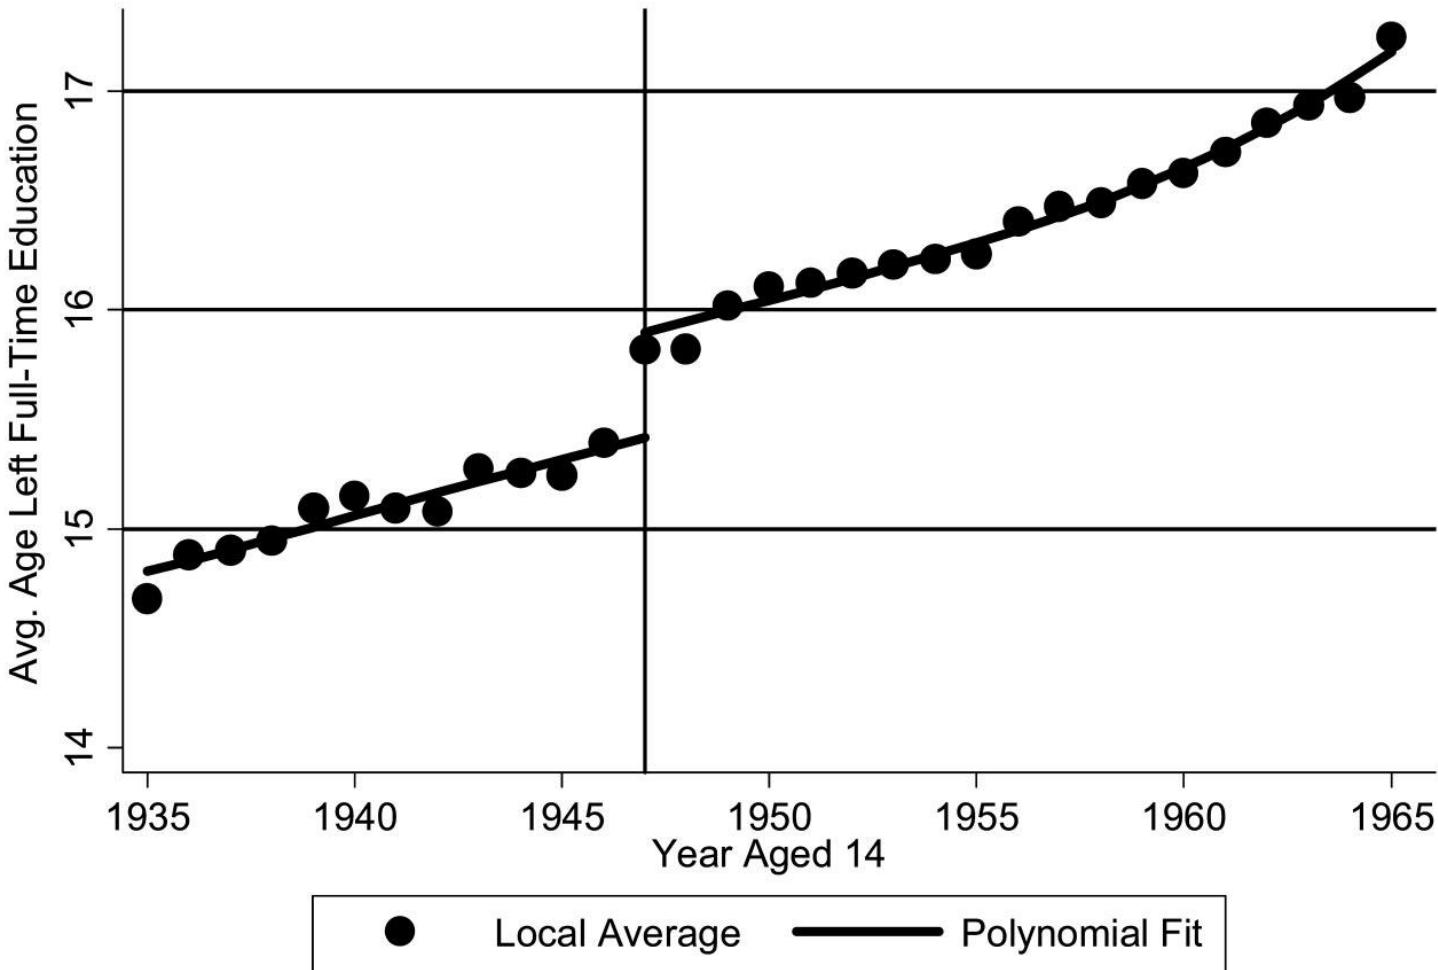
What does a quasi-experiment look like?

Fraction left full-time education by year aged 14 and 15 (Great Britain)



The lower line shows the proportion of British-born adults aged 32 to 64 from the 1983 to 1998 General Household Surveys who report leaving full-time education at or before age 14 from 1935 to 1965. The upper line shows the same, but for age 15. The minimum school-leaving age in Great Britain changed in 1947 from 14 to 15 [Oreopoulos (2006)]

Average schooling increases by exactly half a year between the cohorts that were age 14 in 1946 and in 1948.



Quasi Experiments: Education and Mortality

- Changes in education due to national or state-level changes in laws regarding compulsory schooling.
- Differ from the usual approach by attempting to focus on plausibly random **changes** in education, rather than comparing those **achieving** high vs. low education.
- Findings are heterogeneous, in contrast to much of the evidence from observational studies:
 - USA (Lleras-Muney, 2005): IV (yes), RD (no)
 - UK (Oreopoulos 2008, Clark 2010): RD (no)
 - France (Albouy 2009): RD (no)
 - Also positive and negative evidence for other health outcomes in Denmark, Sweden, Germany, Italy, Netherlands
- Importance of explicitly trying to mimic an RCT for education



Contents lists available at [ScienceDirect](#)

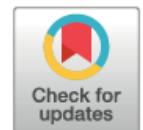
Social Science & Medicine

journal homepage: www.elsevier.com/locate/socscimed



Review article

How and why studies disagree about the effects of education on health: A systematic review and meta-analysis of studies of compulsory schooling laws



Rita Hamad^{a,*}, Holly Elser^{b,c}, Duy C. Tran^c, David H. Rehkopp^c, Steven N. Goodman^c

^a University of California San Francisco, Philip R. Lee Institute for Health Policy Studies, Department of Family & Community Medicine, 995 Potrero Avenue, Building 80, Ward 83, San Francisco, CA, 94110, USA

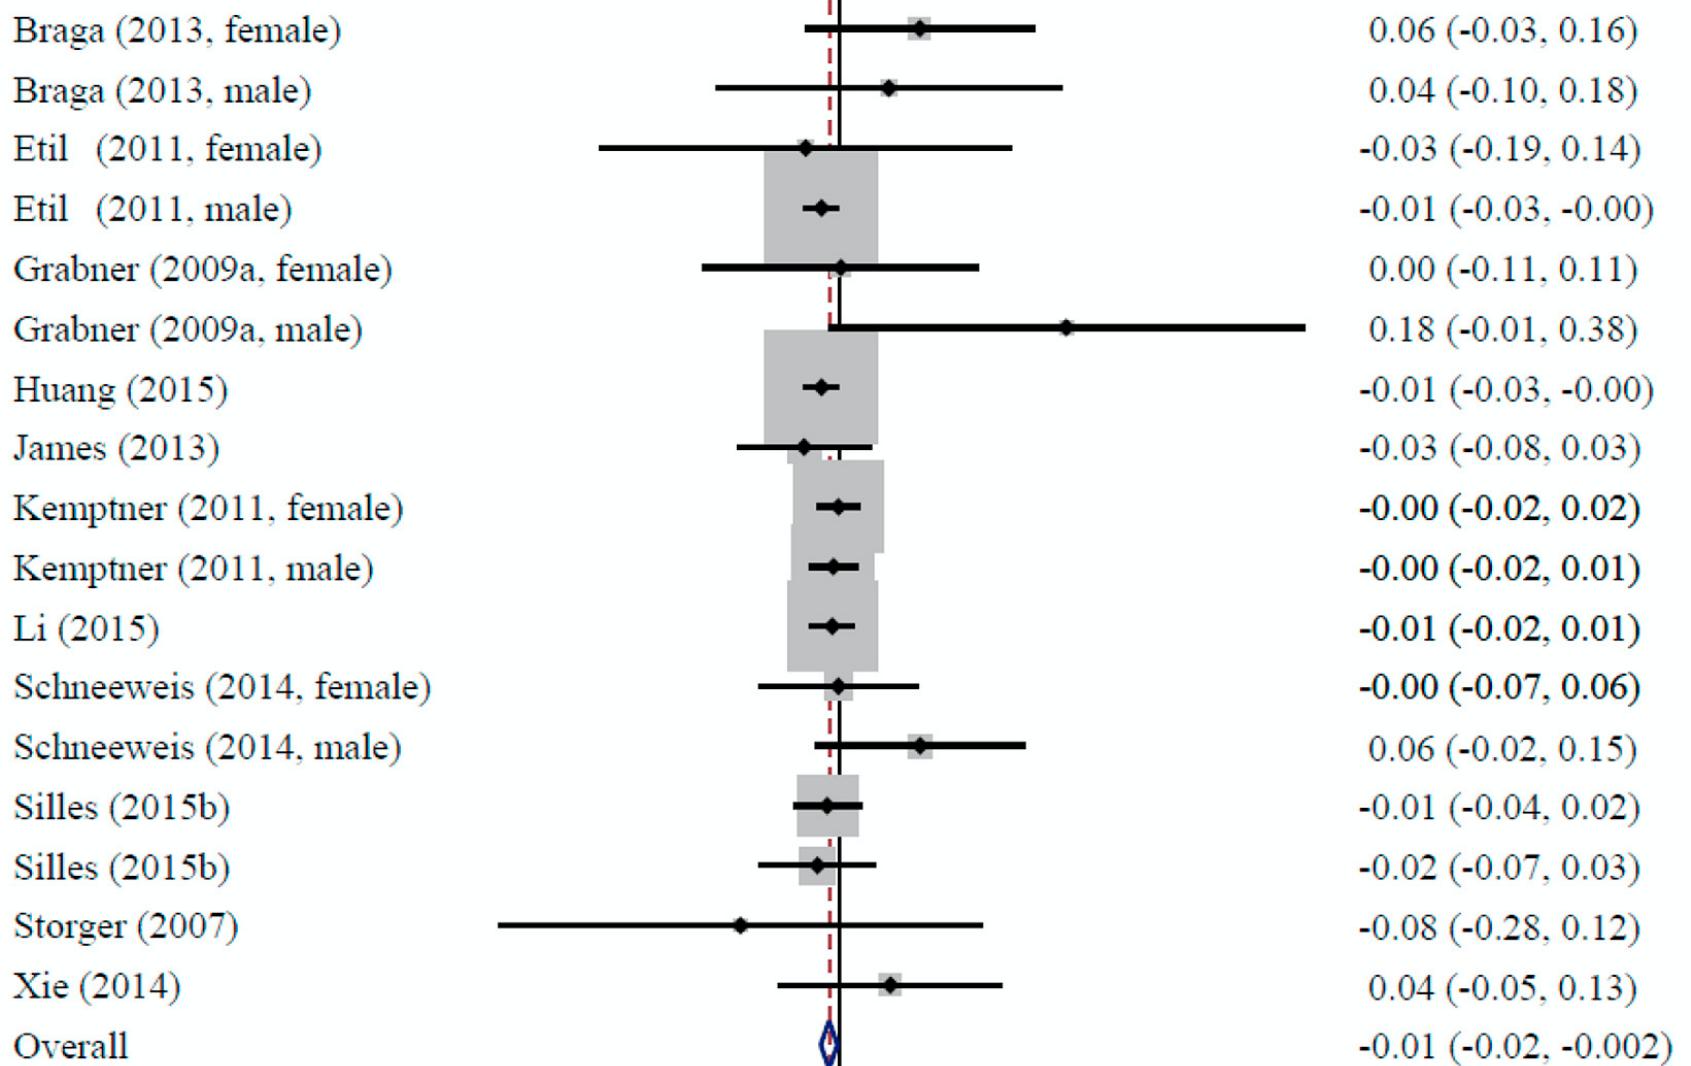
^b University of California Berkeley, School of Public Health, Division of Epidemiology, Berkeley, CA, USA

^c Stanford University, School of Medicine, Stanford, CA, USA

Panel B. Smoking

Study

Effect Size (95% CI)



Regression Discontinuity

RD: Basic Idea

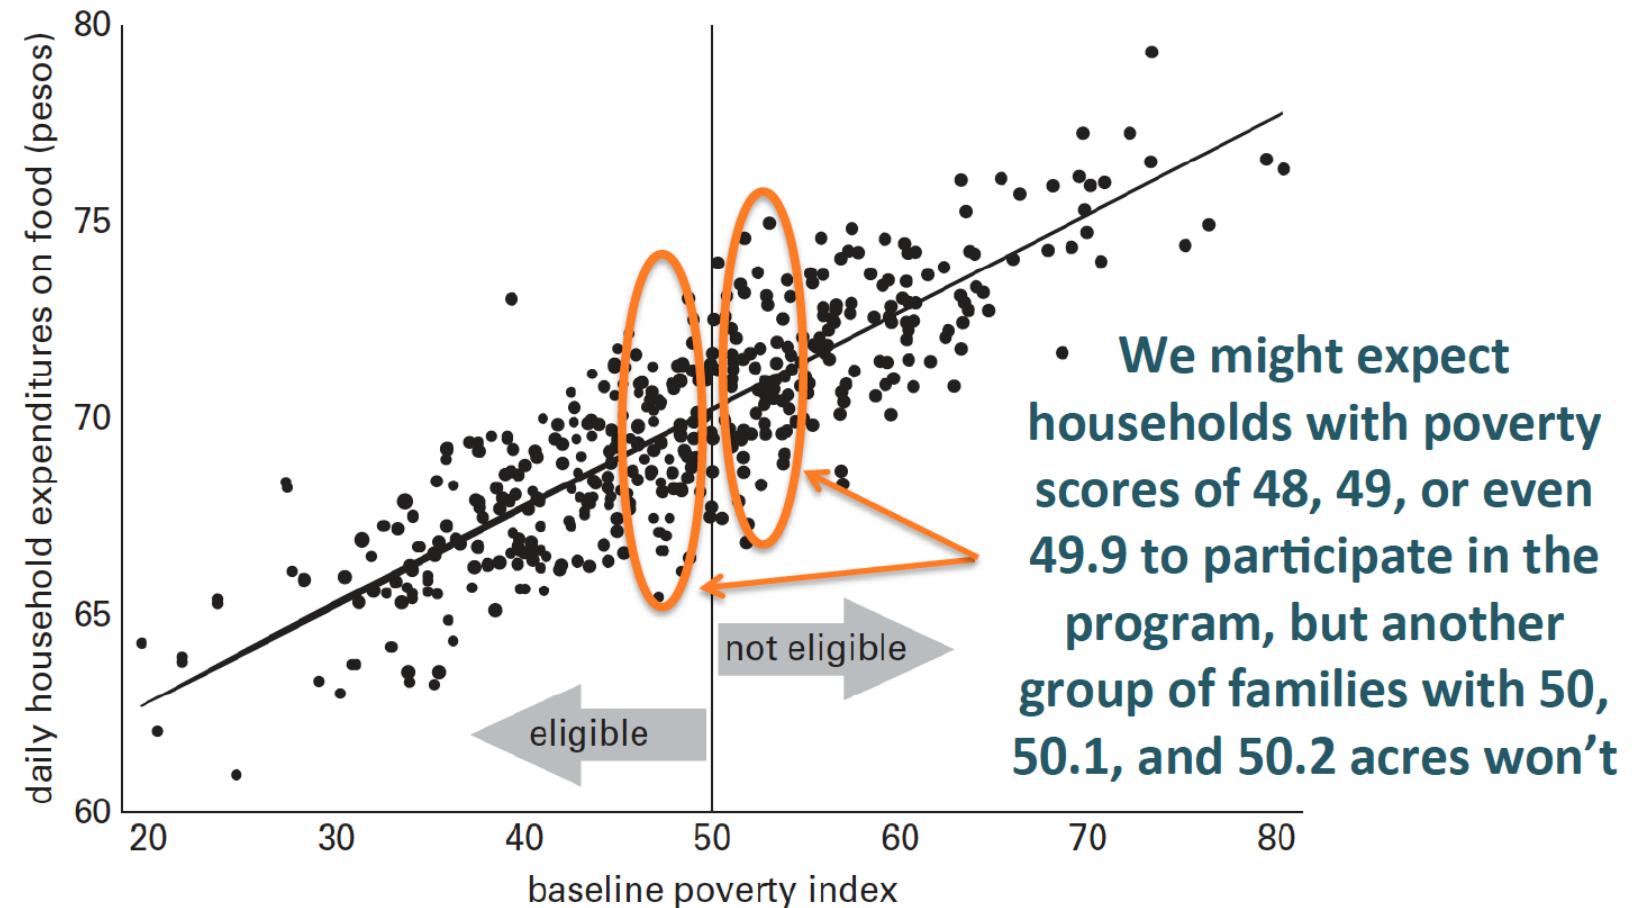
- Take advantage of arbitrary thresholds that sometimes assign treatment to individuals.
- When an administrative or rule-based cutoff in a continuous variable (present in your data) predicts treatment assignment, being on one side or the other of this cutoff determines, or predicts, treatment received.
- The continuous variable is called the "assignment" or "forcing" variable.
- Groups just on either side of the threshold considered "as good as randomly" assigned to treatment.

RD: Motivating example

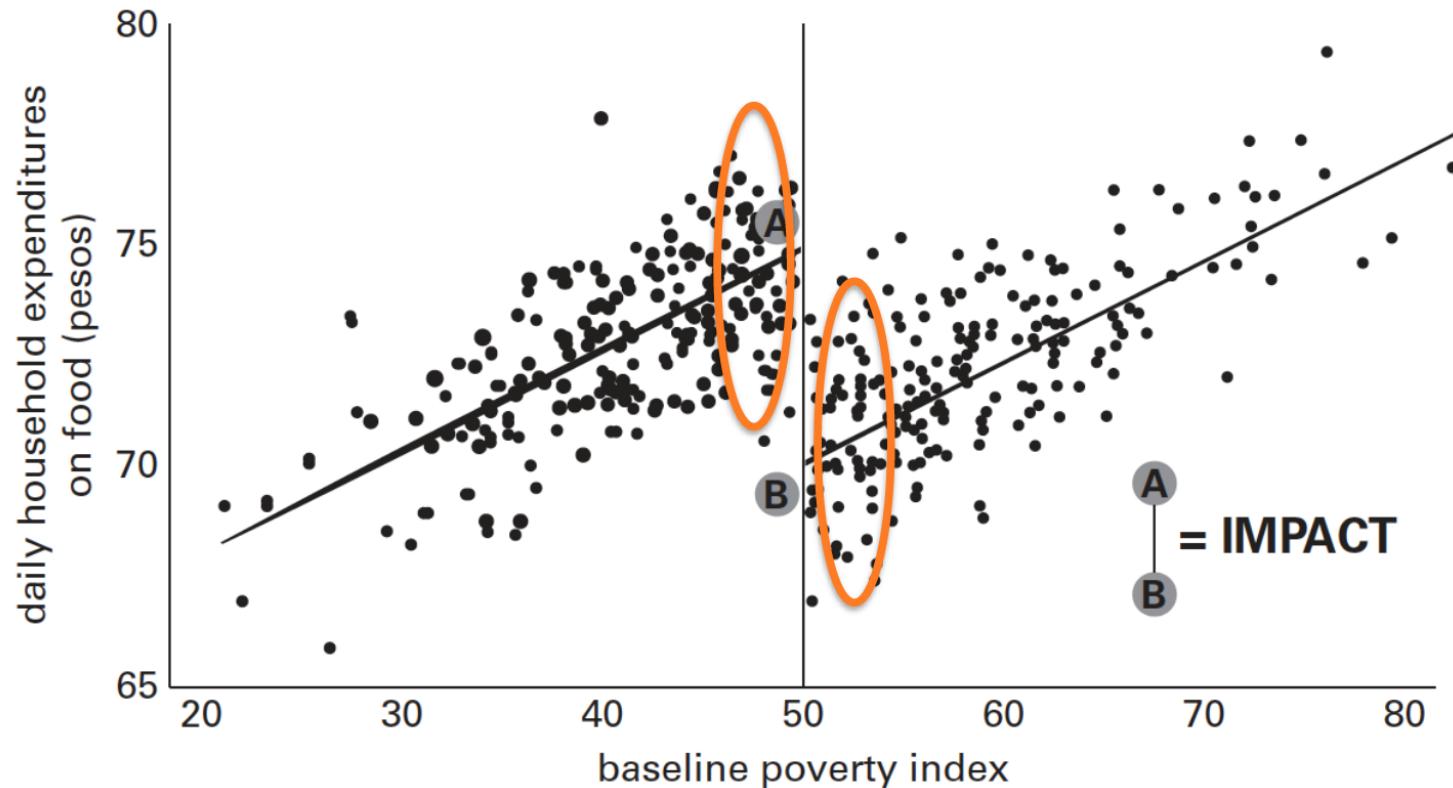
- Suppose we want to estimate the impact of a cash transfer program on daily food expenditure of poor households.
- Poverty is measured by a continuous score between 0 and 100 that is used to rank households from poorest to richest.
- Poverty is the assignment variable, Z , that determines eligibility for the cash transfer program.
- The outcome of interest, daily food expenditure, is denoted by Y .

At baseline, you might expect poorer households to spend less on food, on average, than richer ones, which might look like this

Under the program's rules, only households with a poverty score, Z , below 50 are eligible for the cash payment

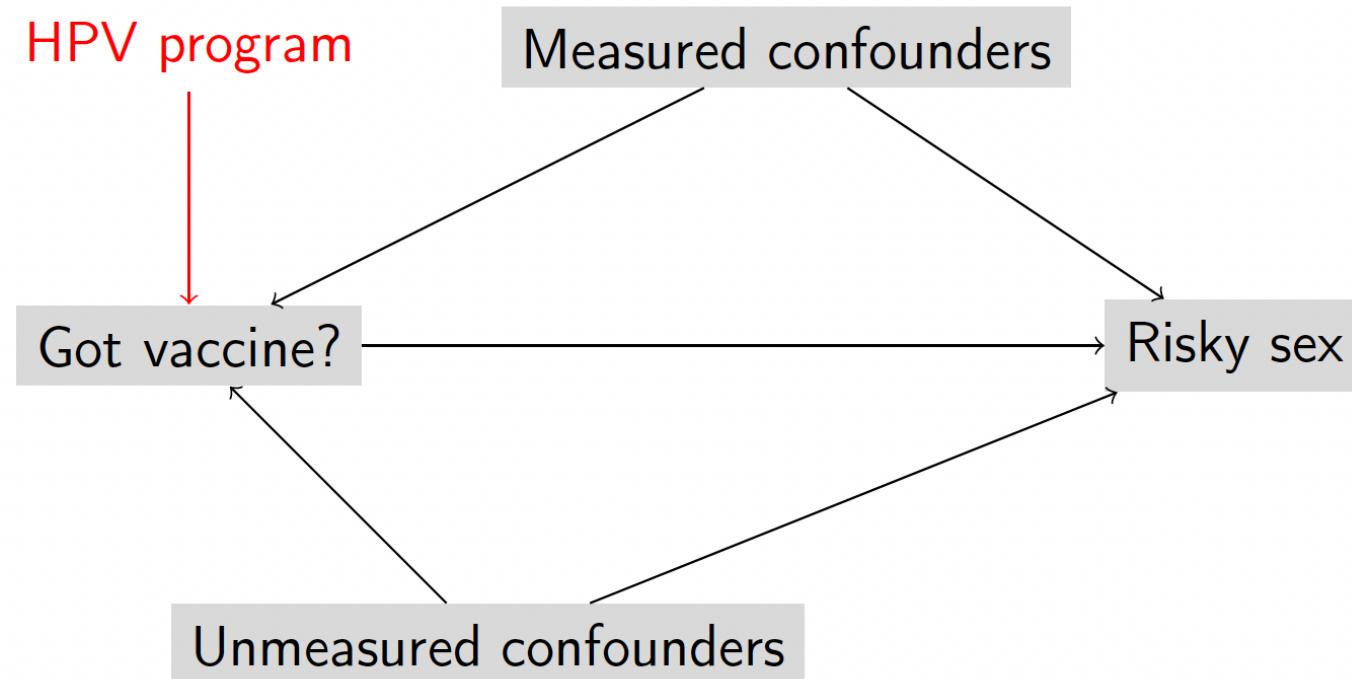


As we approach the cutoff value from above and below, the individuals in both groups become more and more alike, on both measured and unobserved characteristics---in a small area around the threshold, the only difference is in treatment assignment



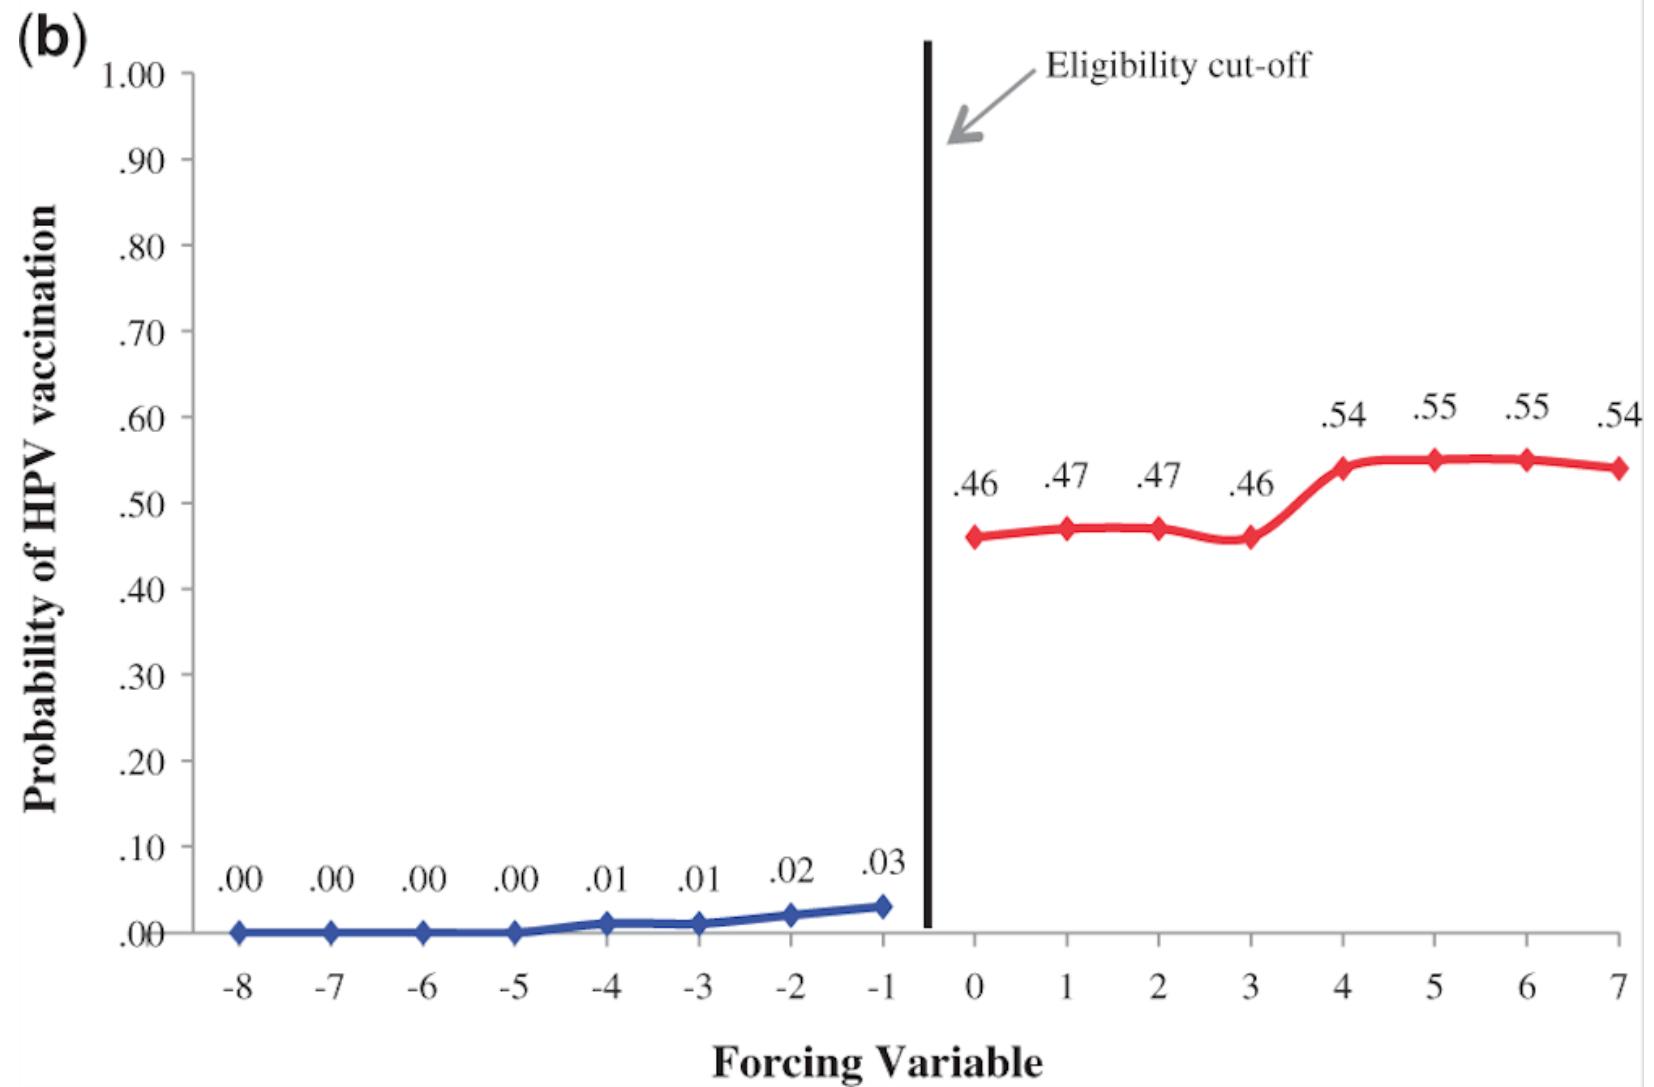
Applied example: HPV vaccine and sexual behaviors

- Does getting the HPV vaccine affect sexual behaviors?
- Vaccine policy: predicts vaccine receipt but (**we assume**) .red[not] associated with anything else [mimicking random assignment].



Does the cutoff predict treatment?

- Girls "assigned" to HPV program by quarter of birth.
- $\text{Pr}(\text{vaccine})$ jumps discontinuously at cutoff



What does a credible natural experiment look like?

Table 1: Baseline characteristics of the eligibility groups in the study cohort

Characteristic	Program eligibility group; % of eligibility group*		Characteristic	Program eligibility group; % of eligibility group*	
	Ineligible (n = 131 781)	Eligible (n = 128 712)		Ineligible (n = 131 781)	Eligible (n = 128 712)
Sociodemographic‡				Health services use**††	
Age, yr, mean ± SD	13.17 ± 0.28	13.17 ± 0.28	Hospital admission		
Birth quarter			0	98.0	98.2
Jan.–Mar.	24.3	24.2	≥ 1	2.0	1.8
Apr.–June	26.1	26.1	LOS, d, mean ± SD	7.4 ± 15.6	8.0 ± 18.2
July–Sept.	25.7	25.8	Same-day surgery		
Oct.–Dec.	23.9	23.9	0	97.7	97.8
Residency			≥ 1	2.4	2.2
Urban	85.3	85.8	Emergency department visits		
Rural	14.0	13.5	0	70.7	71.1
Missing‡	0.7	0.6	1	18.1	17.8
Income quintile			≥ 2	11.2	11.1
1 (lowest)	16.6	15.0	Outpatient visits		
2	18.4	17.8	0 or 1	22.6	22.8
3	20.6	21.1	2–5	27.4	26.9
4	22.0	23.1	6–12	25.1	24.5
5 (highest)	21.4	22.1	≥ 13	25.0	25.8

Smith et al. (2015)

Note little impact of adjustment

Table 3: Effect of quadrivalent human papillomavirus vaccination on clinical indicators of sexual behaviour*

Outcome	No. of excess cases per 1000 girls (95% CI)	RR (95% CI)	Adjusted† RR (95% CI)
Effect of vaccine			
Composite outcome	-0.61 (-10.71 to 9.49)	0.96 (0.81 to 1.14)	0.98 (0.84 to 1.14)
Pregnancy	0.70 (-7.57 to 8.97)	0.99 (0.79 to 1.23)	1.00 (0.83 to 1.21)
STIs	-4.92 (-11.49 to 1.65)	0.81 (0.62 to 1.05)	0.81 (0.63 to 1.04)
Effect of program			
Composite outcome	-0.25 (-4.35 to 3.85)	0.99 (0.93 to 1.06)	1.00 (0.93 to 1.07)
Pregnancy	0.29 (-3.07 to 3.64)	1.00 (0.92 to 1.09)	1.01 (0.93 to 1.10)
STIs	-2.00 (-4.67 to 0.67)	0.92 (0.83 to 1.03)	0.92 (0.83 to 1.03)

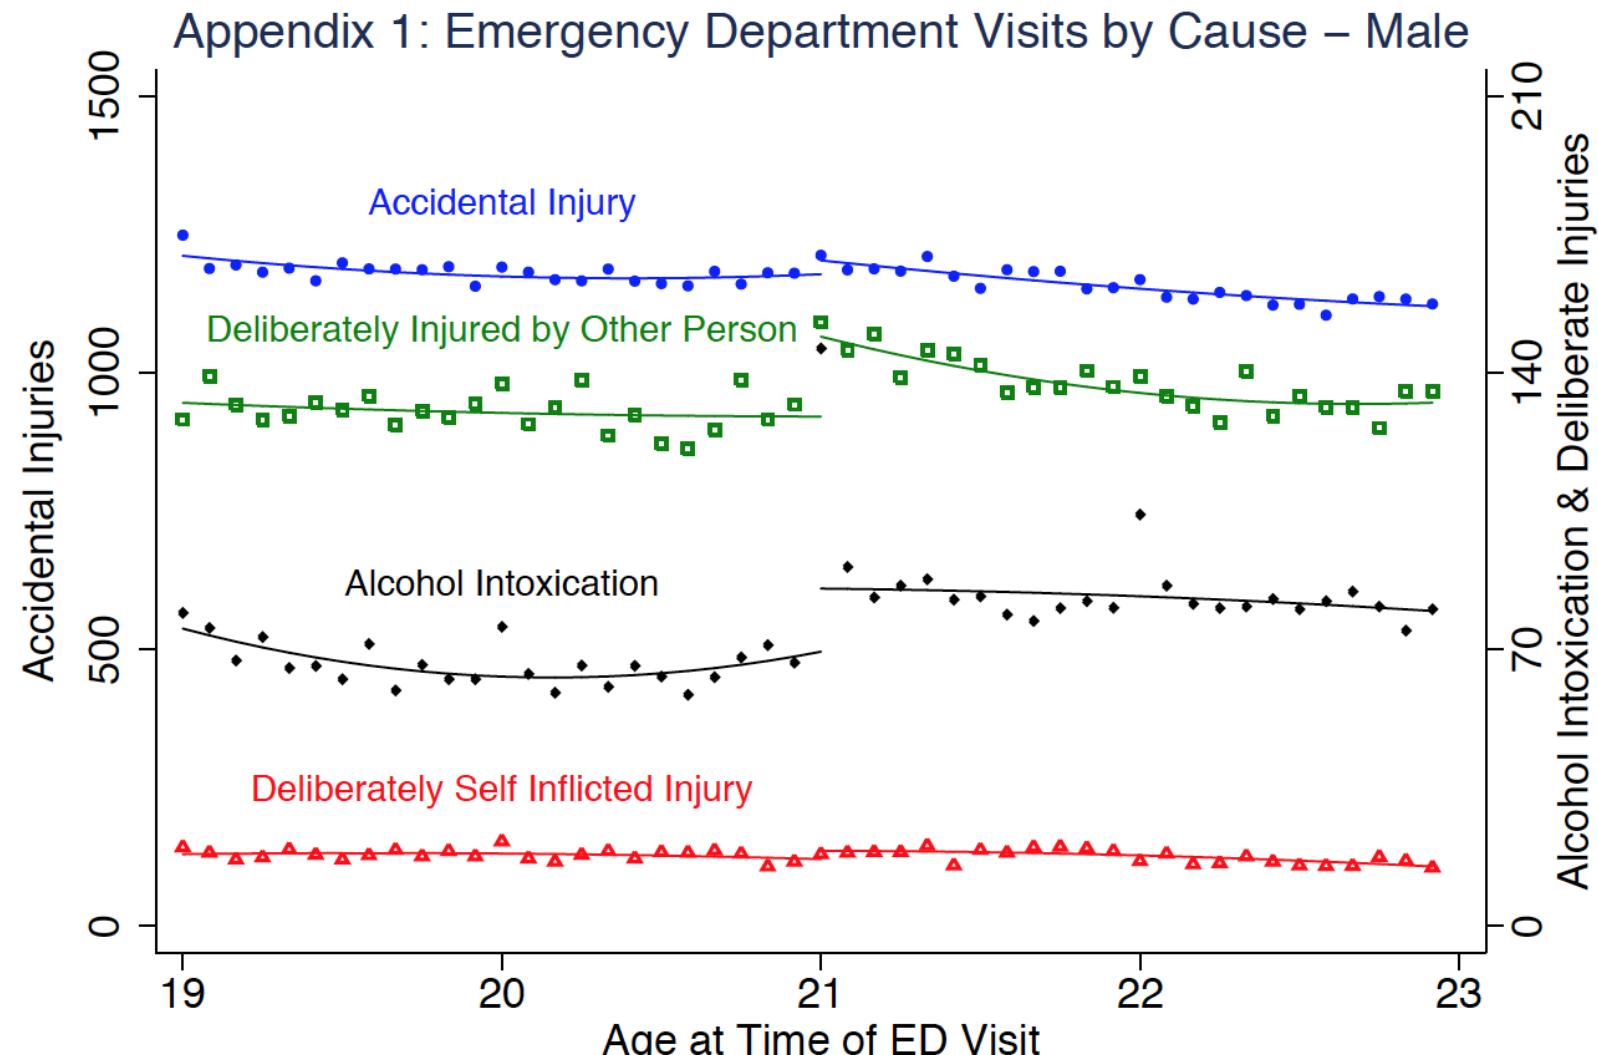
Note: CI = confidence interval, RR = relative risk, STIs = sexually transmitted infections.

*To address the effect of birth timing that we observed, we used the entire bandwidth of data (i.e., all observations in the 1992 to 1995 birth cohorts) and included birth quarter as a covariate in the model. In all analyses, the birth cohorts closest to the cut-off (1993 and 1994) were weighted twice as heavily as those furthest from the cut-off (1992 and 1995).

†In this sensitivity analysis, we adjusted for neighbourhood income quintile, hepatitis B vaccination and history of sexual health-related indicator, as well as for birth quarter.

Another
recent
example: US
drinking age

Minimum legal
drinking age and
non-fatal injuries



Carpenter (2017)

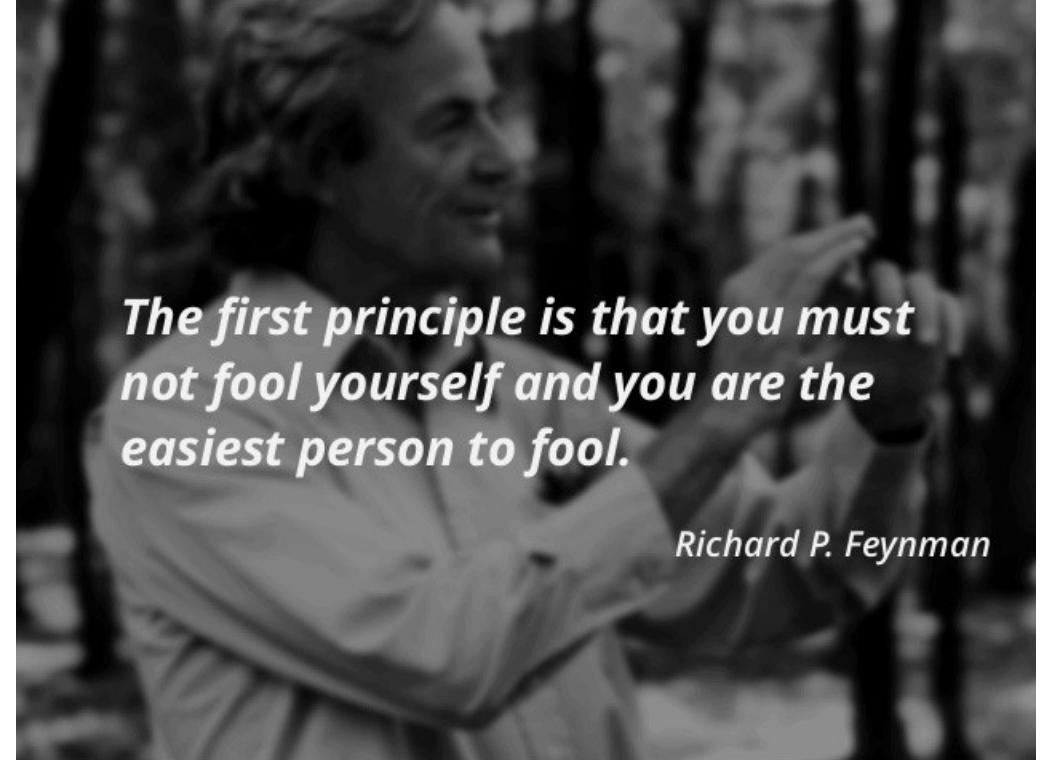
Note: The points are ED visit rates per 10,000 and the fitted lines are from a second order quadratic polynomial in age estimated separately on either side of the threshold.

Issues related to generalizability

- RD estimates local average impacts around the eligibility cutoff where treated and control units are most similar and results cannot be generalized to units whose scores are further away from the cutoff (unless we assume treatment heterogeneity).
- If the goal is to answer whether the program should exist or not, then RD is likely not the appropriate methodology.
- However, if the question is whether the program should be cut or expanded at the margin, then it produces the local estimate of interest to inform this policy decision

Be careful, and skeptical

- Correlations between social factors and health are easy to find.
- They do not necessarily reflect **causal** relationships.
- Need to search hard for alternative explanations.
- Important to consider the strength of evidence in considering interventions.



The first principle is that you must not fool yourself and you are the easiest person to fool.

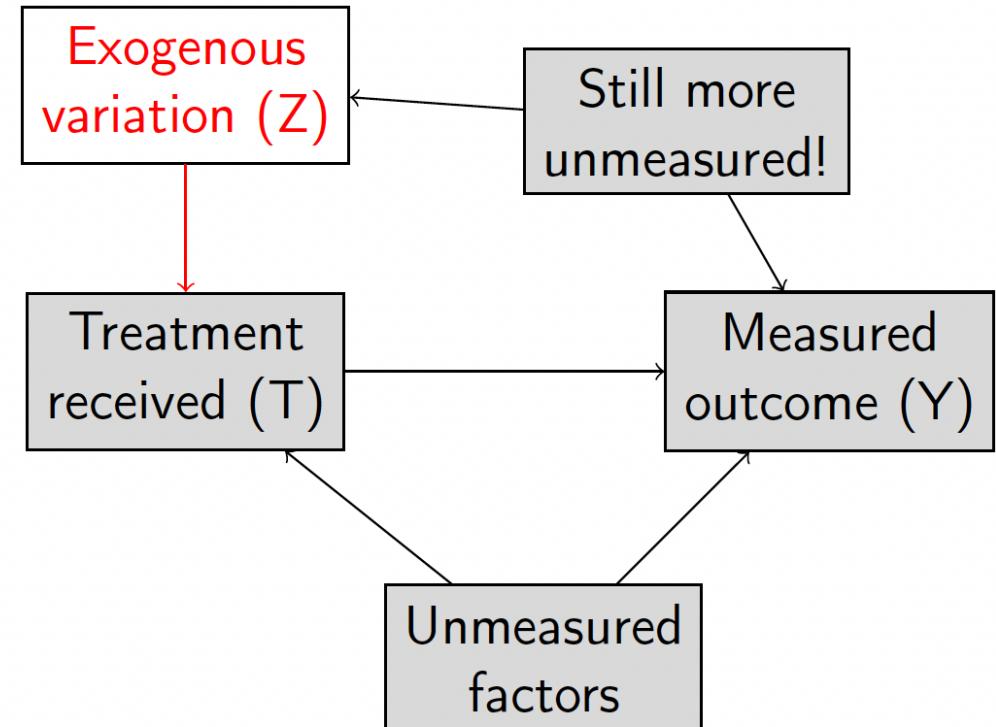
Richard P. Feynman

Are natural experiments always more credible?

- Not necessarily, but probably.
- Key is "as-if" randomization of treatment:
 - If this is credible, it is a much stronger **design** than most observational studies.
 - Should eliminate self-selection into exposure groups.
- Allows for simple, transparent analysis of average differences between groups.
- Allows us to rely on weaker assumptions.

Assumptions still matter!

- Quasi-experimental studies are still observational.
- Most credible if they create unconditional randomized treatment groups (e.g., lottery).
- Credibility is continuous, not binary.
- I worry about the cognitive impact of the "quasi-experimental" label.



Potential drawbacks of quasi-experimental approaches

- How good is "as-if" random? (need "shoe-leather")
- Credibility of additional (modeling) assumptions.
- Relevance of the intervention.
- Relevance of population.

But potential benefits!

However, this year's Laureates have shown that it is possible to answer these and similar questions using natural experiments. The key is to use situations in which chance events or policy changes result in groups of people being treated differently, in a way that resembles clinical trials in medicine.

Royal Swedish Academy of Sciences (2021)



Back to basics: assumptions and costs

- Major benefit of randomized evaluations are that few assumptions are needed to estimate a causal effect.
- Necessary assumptions can often be checked.
- Non-randomization means more assumptions, more possibility for assumptions to be violated.
- Should lead us to spend lots of time trying to test the credibility of these assumptions.
 - How good is "as-if random"?
 - Are there compelling non-causal alternative explanations for the observed results?
- All non-randomized designs are not created equal.